



# Pipelines in Unix Systems

An overview of pipeline processing


Author: Rodrigo Gonzalez, PhD



# Overview of UNIX Pipes

- © Pipes are used for IPC (Inter-Process Communication) in UNIX, allowing data flow from one process to another.
- © Example usage in shell commands: ``$ ls | wc -l`` to count the number of files in a directory.
- © The shell creates two processes for ``ls`` and ``wc`` commands, utilizing a pipe for data transfer between them.

# Creating and Using Pipes



**Figure 44-1:** Using a pipe to connect two processes

# Characteristics of Pipes

- © Pipes operate as byte streams, lacking the concept of messages or boundaries.
- © Data is read in the same order it was written; random access is not possible.
- © Pipes are inherently unidirectional, though bidirectional alternatives exist (e.g., UNIX domain stream socket pairs).

# Reading and Writing to Pipes

- © Attempts to read from an empty pipe block until data is available.
- © Writes up to `PIPE_BUF` bytes are **atomic**, preventing data intermingling from multiple processes.

# Capacity and Blocking Behavior of Pipes

- © Pipes have a limited capacity; exceeding it causes writes to block until space is freed.
- © Linux allows modification of pipe capacity, affecting how data is buffered and transferred.

# Creating a Pipe

- © System Call: `pipe()` creates a pipe and returns two file descriptors, one for reading and one for writing.
- © Data Flow: Data written to the write end is immediately available at the read end.


```
#include <unistd.h>
```

```
int pipe(int filedes[2]);
```

Returns 0 on success, or -1 on error

# Creating a Pipe

- © System Call: `pipe()` creates a pipe and returns two file descriptors, one for reading and one for writing.



**Figure 44-2:** Process file descriptors after creating a pipe

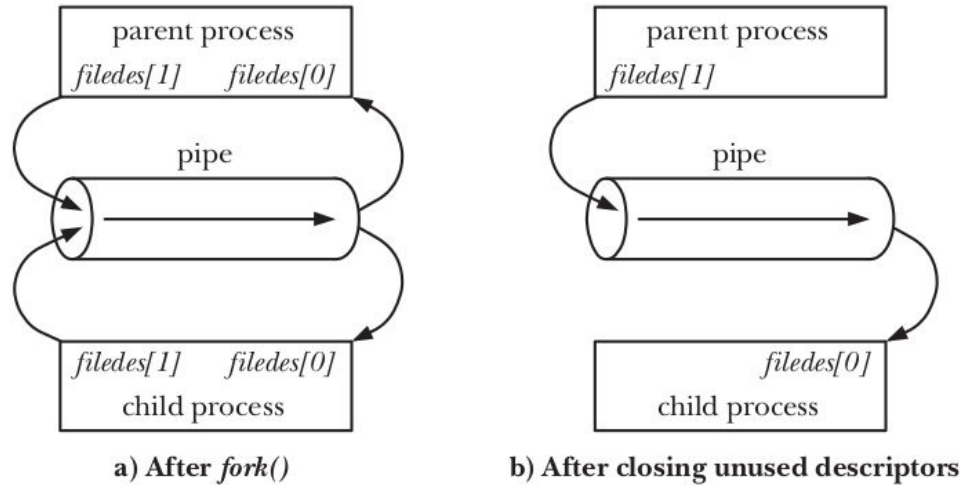


# Pipes in Process Communication

- © Parent-Child Communication: Pipes are typically used to connect parent and child processes for data exchange.
- © Operational Detail: The process not using one end of the pipe (read or write) should close that end.

# Pipes in Process Communication

- © Parent-Child Communication: Pipes are typically used to connect parent and child processes for data exchange.



**Figure 44-3:** Setting up a pipe to transfer data from a parent to a child

# Practical Considerations

- © The parent must close the write end it owns to ensure it properly detects EOF.
- © This method can be adapted for complex synchronization schemes using multiple pipes or additional messaging.

# Introduction to Process Synchronization Using Pipes

- © Overview: Pipes not only facilitate data transfer but also serve as a synchronization mechanism.
- © Context: Race conditions in parent-child processes can be mitigated using pipes for synchronization.

# Synchronization Mechanism

- ◎ The parent process creates a pipe prior to forking child processes.
- ◎ Each child process performs its designated action and then closes its write end of the pipe.
- ◎ The parent process, after closing its write end, reads from the pipe to wait for all children to complete.

# Advantages of Pipe-based Synchronization

- © Synchronization without busy waiting: Parent process waits on read, consuming minimal resources.
- © Scalability: Easily extendable to synchronize multiple child processes with a single parent.

# Implementing Synchronization

- © Closing the write end of the pipe in child processes triggers EOF in the parent process's read operation.
- © The EOF signals to the parent process that all child processes have completed their actions.

# Comparing Synchronization Techniques

- © Pipes vs. Signals: Pipes provide a deterministic way to synchronize multiple related processes, unlike signals which cannot be queued.
- © Flexibility: Pipe-based synchronization allows for straightforward, expandable coordination among processes.



# Closing Remarks on Pipes

- © Flexibility: Pipes offer a simple yet powerful way to handle IPC with immediate data availability.
- © Best Practices: Proper management of file descriptors and understanding unidirectional data flow are key to leveraging pipes.